

PALFA Pipeline2.0 Documentation

Patrick Lazarus, Alex Samoilov, ...
plazar@physics.mcgill.ca, alex@sequencefactory.com

April 20, 2011

Abstract

1 Installation

1.1 Dependencies

The PALFA pipeline has various dependencies. All of the following packages are required to run the pipeline.

- PRESTO (a recent version is required. The new executable `fitsdelcol` must be included - <https://github.com/scottransom/presto>)
- `psrfits_utils` (be sure to pull Kevin Stovall's version of `psrfits_utils`. It contains the merging code required to analyse PALFA Mock spectrometer data. https://github.com/kstovall/psrfits_utils.git)
- `numpy` (<http://numpy.scipy.org/>)
- `M2Crypto` (<http://pypi.python.org/pypi/M2Crypto>)
- `PBSQuery` (if PBS is being used - <http://subtrac.sara.nl/oss/pbs-python/wiki/TorqueInst>)
- `pyfits` (http://www.stsci.edu/resources/software_hardware/pyfits)
- `pyodbc` (<http://pypi.python.org/pypi/pyodbc/>)

- ImageMagick (on worker nodes - <http://www.imagemagick.org/script/index.php>)
- FreeTDS (required by pyodbc to connect to common-DB - <http://www.freetds.org/>)
- UnixODBC (required by pyodbc to connect to common-DB - <http://www.unixodbc.org/>)
- prettytable (<http://code.google.com/p/prettytable/>)

*NOTE: Some of these dependencies have requirements of their own. For example, PRESTO requires *cfitsio*, *TEMPO*, *numpy*, *scipy*, etc. Be sure to follow the installation instructions for each packages.*

1.2 Getting Started

Here we will present basic step-by-step instructions for setting up the pipeline.

Step 1 Create a directory where you want the pipeline to be installed.

Step 2 Download the pipeline source files. This should be done by cloning the git repository on github.com: <https://github.com/plazar/pipeline2.0>.

```
$ git clone git://github.com/plazar/pipeline2.0.git
```

NOTE: This will create a sub-directory called “pipeline2.0”.

Step 3 Add the pipeline’s lib/python directory to your PYTHONPATH environment variable.

Step 4 Create configuration files using the examples provided. Modify the settings so they are appropriate for your system. It is possible to check the sanity of your configurations by calling a configuration file with the python interpreter. An error message will be displayed if a configuration type or value is invalid.

Step 5 Configure ODBC and FreeTDS. This can be done by creating the appropriate files (`.odbc.ini` and `.freetds.conf`, respectively) in your home area. See Section 3.1 for details.

Step 6 Set up the database. This is done using the `create_database.py` script.

Step 7 Perform tests. In the `test/` directory there are some tests that can be run to see if you can connect to the common DB (`commondb_test.py`), e-mail system (`mailer_test.py`), and FTP server (`cornellftp_test.py`). To perform the tests call each script with the python interpreter. Additional information about the tests can be found in Section 3.2.

Step 8 Start the downloader using `StartDownloader.py`.

NOTE: The downloader does not need to be run on the same computer as the job pooler. However, the directory where the downloader saves files must be accessible from the computer where the job pooler is being run.

Alternatively, it is possible to add files to the job-tracker database without using the downloader. To do this use `add_files.py`. Be sure to set the `delete_rawdata` configuration to `False` if you do not want raw data files to be deleted when the pipeline no longer requires them.

Step 8 Start the job pooler using `StartJobPool.py`. The job pooler will start submitting jobs when data files are finished downloading.

Step 9 Start the uploader using `StartJobUploader.py`. The uploader will upload results to the common database when jobs are successfully processed.

2 The Pipeline

An overview of the pipeline.

2.1 Features

- Sanity check of configurations
- Dynamic zapping
- Automatic download of data files
- File size checks when downloading
- Notification of errors by email

- Automatic (configurable) retry of failed jobs
- Results are uploaded as a single transaction, so no partial uploads when errors occur
- Automatic (configurable) deletion of data files (upon success or terminal failure)

2.2 Components

2.2.1 Job-tracking Database

Each job run through the pipeline passes through multiple stages before it is completed. To track the state of each job, and log its history a SQLite3 database is used. The database is simply a (readable/writable) binary file formatted according to the SQLite3 format. The database file can be accessed using SQL statements using python's sqlite3 module, which is part of python's standard library, as of python version 2.5.

Relevant links:

- <http://docs.python.org/library/sqlite3.html>
- <http://www.sqlite.org/>

The job-tracker database governs all aspects of the pipeline. It maintains lists of requests for data, downloads, jobs, files associated with each job, processing attempts, and upload attempts.

Details for each table in the database are presented in Section 3.6.

2.2.2 Downloader

2.2.3 Job Pooler

2.2.4 Uploader

Headers

Candidates

Diagnostics Diagnostics stored in the common database are split into two types: Numeric values and binary data. Some of the “binary data” diagnostics are actually plain text files.

The following per-beam diagnostics are uploaded to the common database:

- *RFI mask percentage* (Numeric value)
Percentage of data masked due to RFI.
- *Num candd folded* (Numeric Value)
The number of candidates folded.
- *Num candd produced* (Numeric Value)
The total number of candidates produced, including those with sigma lower than the folding threshold.
- *Min sigma folded* (Numeric value)
The smallest sigma value of all folded candidates from this beam.
- *Num candd above threshold* (Numeric value)
The number of candidates produced (but not necessarily folded) that are above the desired sigma threshold.
- *RFI find png* (Binary data)
Output image produced by rfind in png format.
- *Accelcandd list* (Binary data)
The combined and sifted list of candidates produced by accelsearch.
- *Zaplist used* (Binary data)
The list of frequencies and ranges zapped from the power spectrum before searching this beam.

2.3 Objects Tracked

Restores

Downloads

Jobs

Uploads

3 Reference

3.1 Database Configuration

UnixODBC and FreeTDS need to be properly configured to be able to connect to the common database. The following configuration files are required (at McGill). These files can be placed in your home area. It is also possible to put these configurations in a system configuration directory so all users will have access to them automatically.

`.odbc.ini`

```
[ODBC Data Sources]
FreeTDSdsn = FreeTDS
MySQLdsn   = MySQL

[FreeTDSdsn]
Driver      = /usr/lib/odbc/libtdsodbc.so
Servername  = Cornell
Description = FreeTDS
Database    = palfa-common-copy

[MySQLdsn]
Driver      = /usr/lib/odbc/libmyodbc.so
Description = MySQL

[Default]
Driver      = /usr/lib/odbc/libtdsodbc.so
```

`.freetds.conf`

```
# $Id: freetds.conf,v 1.12 2007/12/25 06:02:36 jklowden Exp $
#
# This file is installed by FreeTDS if no file by the same
# name is found in the installation directory.
```

```

#
# For information about the layout of this file and its settings,
# see the freetds.conf manpage "man freetds.conf".

# Global settings are overridden by those in a database
# server specific section
[global]
    # TDS protocol version
;    tds version = 4.2

    # Whether to write a TDSDUMP file for diagnostic purposes
    # (setting this to /tmp is insecure on a multi-user system)
;    dump file = /tmp/freetds.log
;    debug flags = 0xffff

    # Command and connection timeouts
;    timeout = 10
;    connect timeout = 10

    # If you get out-of-memory errors, it may mean that your client
    # is trying to allocate a huge buffer for a TEXT field.
    # Try setting 'text size' to a more reasonable limit
    text size = 16777215

# A typical Sybase server
[egServer50]
    host = symachine.domain.com
    port = 5000
    tds version = 5.0

# A typical Microsoft server
[egServer70]
    host = ntmachine.domain.com
    port = 1433
    tds version = 7.0

[Corne11]

```

```
host = arecibosql.tc.cornell.edu
port = 1433
```

3.2 Tests

The pipeline contains several tests that can be used to confirm certain tricky aspects of the pipeline are configured properly. These tests are located in `tests/`.

commondb_test.py This script tests that connections to the common database. It also verifies that it is possible to read from the database, as well as insert entries into the database.

cornellftp_test.py This test confirms that it is possible to connect to the FTP server at Cornell, access information about files available, and retrieve files.

install_test.py This script attempts to import all python modules necessary by the pipeline. If any modules were not successfully imported they are reported to the user.

mailer_test.py This script attempts to send an email using the SMTP server defined in the configuration file `email.py`.

restore_test.py This script verifies that the pipeline can connect to the data-restore web services, make requests and check when data are available.

3.3 Executables

The pipeline has a series of executables located in the `'bin/'` directory. There are two categories of executables: background scripts, and utilities.

3.3.1 Background scripts

The following background scripts are used to run the Downloader, Job Pooler and Uploader.

- **StartDownloader.py**
A script to start the downloader background process.
usage: **StartDownloader.py**
- **StartJobPool.py**
A script to start the job pooler background process.
usage: **StartJobPool.py**
- **StartJobUploader.py**
A script to start the results uploader background process.
usage: **StartJobUploader.py**

3.3.2 Utilities

The following utility scripts are used to interact with the job-tracker database in a safe way. This includes creating the database, editing the entries in the database, and compiling and displaying information from the database entries.

- **add_files.py**
A utility to add files from a given directory to be tracked. This is useful for tracking files that are copied over from an external disk, as opposed to downloaded by the Downloader. Note that added files will be deleted just like downloaded files.
usage: **add_files.py** DIRECTORY
- **stop_processing_jobs.py**
Stop a job running in the queue. There are two ways to stop jobs: 1) Failing the job (i.e. the submission counts towards the job's number of retries, and 2) Removing the job (the submission doesn't count towards retries). Both possibilities are done safely, with respect to the job-tracker DB. The default is to remove the jobs (not fail).
usage: **stop_processing_jobs.py** [OPTIONS] QUEUE_ID [QUEUE_ID ...]
options:
-h, --help Show help message and exit.
-f, --fail Remove jobs from the queue and mark them as 'failed' in

the job-tracker database. (Default: Remove jobs and don't mark them as 'failed').

- **create_database.py**

Create the sqlite job-tracker database. Note that the script will do nothing if the database already exists. The location of the database is taken from the database configuration file.

usage: create_database.py

- **kill_jobs.py**

Kill a job. That is set its status as 'terminal_failure', and clean up its datafiles (if applicable).

usage: kill_jobs.py ID [ID ...]

options:

-h, --help Show help message and exit.

-f FILES, --file=FILES File belonging to a job that should be killed.

-i JOBIDS, --id=JOBIDS ID number of a job that should be killed.

3.4 Generic Queue Manager Interface

The Pipeline interacts with the system's queue manager using an abstract interface. This generic interface can be implemented with the specific queue manager installed on the system. The system-specific queue manager class must be derived from the `generic_interface.PipelineQueueManager` class. Each of the following methods must be implemented in the derived class:

submit Submits a job to queue manager with files list and output directory. Must return a unique string identifier for the submitted job. Raises `pipeline_utils.PipelineError` if job submission fails.

can_submit Return a boolean value that indicates if jobs can be submitted to the queue.

is_running Given a unique identifier for a job, must return True or False whether the job is running or not, respectively.

delete Given a unique identifier for a job, remove the corresponding job from the queue. Raises `pipeline_utils.PipelineError` if job removal isn't successful.

status Return the status of the queue as a tuple containing the number of jobs running and the number of jobs queued.

had_errors Given a unique identifier for a job, return True, or False whether the job terminated with an error, or not, respectively.

get_errors Given a unique identifier for a job, return the contents of the error log. If there were not errors return an empty string.

3.5 Configurations

The pipeline requires a series of configurations to be set in order to run. The configurations for different aspects of the pipeline are divided into the different files. The configuration files can be found in the pipeline’s `lib/python/config/` directory.

A sanity check of configurations is performed each time a configuration file is imported. If there are errors (e.g. a setting isn’t provided, a directory doesn’t exist, etc.) a message is printed to the terminal, and an exception is raised causing the program to exit. It is possible to run the check the sanity of a configuration file by running it. For example, checking the sanity of `basic.py`:

```
python basic.py
```

If the settings are sane there will be no error message is output. *NOTE: it may be necessary to check the sanity of the processing configurations on the worker nodes.*

Each of required settings in each of the configuration files are listed below:

3.5.1 `basic.py`

Basic pipeline parameters are found in `basic.py`.

institution Institution where the pipeline is being run. This is the value that will get reported in the common database.

pipeline The type of pipeline being run. This should be left as “PRESTO”.

survey The name of the survey. This should be left as “PALFA2.0”.

pipelinedir The path of the base directory containing the pipeline2.0 code.

psrfits_utilsdir The path of the directory containing the installation of `psrfits_utils` that is used by the pipeline.

delete_rawdata A boolean value that determines if raw data is deleted when results for a job are successfully uploaded to the common DB, or if the maximum number of attempts for a job is reached. *Set this value to False if you do not want the local copy of your data deleted.*

coords_table The path to the `PALFA_coords_table.txt` file, which contains correct coordinates for WAPP data files.

log_dir The path of the directory where log files are stored.

qsublog_dir The path where stderr streams from processing jobs are sent.

3.5.2 background.py

General settings for the three background scripts are found in `background.py`

screen_output A boolean value that should be set to True if you want log information copied to the terminal.

jobtracker_db The path to the SQLite job-tracker database.

sleep The number of seconds to sleep between iterations of the background scripts' loops.

3.5.3 commondb.py

The information needed to connect to the common database is in `commondb.py`.

username The common database user name.

password The password used to access the common database.

host The computer through which a connection to the common database will be established. This should be left at "arecibosql.tc.cornell.edu".

3.5.4 download.py

Information necessary to connect to the FTP server at Cornell and use the data-restore web applications are set in `download.py`, so are settings for controlling the behaviour of the downloader.

api_service_url This is the URL of the data-restore web apps. This should be set to “`http://arecibo.tc.cornell.edu/palfadataapi/dataflow.asmx`”.

api_username This is your user name used to connect to the data-restore web apps.

api_password This is the password for connecting to the data-restore web apps.

ftp_host This is the name of the FTP server at Cornell. It should be left at “`arecibo.tc.cornell.edu`”.

ftp_port This is the port to use for FTP. It should be 31001.

ftp_username This is the user name to connect when FTP’ing. Currently all downloaders are using a shared account, so this should have the value “`palfadata`”.

ftp_password This is the password of the FTP account.

temp This is the directory where data are downloaded to.

space_to_use This is the space available for downloaded files, measured in bytes.

numdownloads This is the number of downloads that can be run in parallel. Currently the downloader is not threaded, so this setting is ignored.

numrestores This is the number of requests for data that can be active at any time. Since the downloader is not threaded only 1 request can be processed at a time, so this setting is also ignore.

numretries This is the number of times a download will be attempted before giving up and moving on to another file.

log_file_path The path of the file where the downloader logs information.

3.5.5 email.py

The pipeline can be configured to send emails when failures occur. The relevant settings are in `email.py`.

enabled A boolean value that enables mailing. If False emails will not be sent.

smtp_host Emails are sent using the Simple Mail Transfer Protocol (SMTP). This is the host of the SMTP server to use. If set to `None` the localhost will be used as a SMTP server to send the emails.

smtp_port The emailer connects to the SMTP server using this port. Either 25 or 587 should work, possibly both.

smtp_username The user account to connect to the SMTP server.

smtp_password The password to connect to the SMTP server.

recipient The email address that will be sent notifications.

sender The email address that should appear in the “From” field of the email.

send_on_failures A boolean value that determines if emails should be sent when a job fails.

send_on_terminal_failures A boolean value that determines if emails should be sent on terminal failures (i.e. job will not be re-tried). Note: if “`send_on_failures`” is True emails will be sent on terminal failures regardless of the value of “`send_on_terminal_failures`”.

3.5.6 jobpooler.py

The setting that determine the behaviour of the Job Pooler are in `jobpooler.py`.

base_results_directory The path of where results from processing are stored. Note: the processing of each data set is in a separate subdirectory.

max_jobs_running The maximum number of jobs running at any given time.

max_jobs_queue The maximum number of jobs that are allowed to be queued at any given time. This must be at least 1, otherwise the Job Pooler will not submit jobs for fear of them being queued.

max_attempts The maximum number of times a job is submitted before being abandoned.

queue_manager An instance of a sub-class of `PipelineQueueManager`. This is the interface between the pipeline and the queue manager.

3.5.7 processing.py

Setting required during processing (i.e. on the worker nodes) are provided in `processing.py`.

base_working_directory The directory on the worker node where temporary working and results directories will be placed.

default_zaplist The path of the zaplist to use when no MJD-specific zaplist can be found.

zaplistdir The directory containing MJD-specific zaplists.

3.5.8 searching.py

The search parameters of the pipeline are located in `searching.py`. To ensure searching is done uniformly across all processing sites most of these values shouldn't be modified.

use_subbands A boolean value that determines if we'll dedisperse and fold using subbands. In general, it is a very good idea to use them if there is enough scratch space on the machines that are processing ($\sim 30\text{GB}/\text{beam}$ processed).

fold_rawdata A boolean value that determines if folding should be done from the raw FITS files (as opposed to subbands or dedispersed FITS files).

datatype_flag This is the flag provided to PRESTO programs to indicate the data type. This should be `"-psrfits"`.

rfifind_chunk_time The duration of each chunk in **rfifind**.

singlepulse_threshold The threshold SNR for singlepulse candidate determination.

singlepulse_plot_SNR The threshold SNR for displaying singlepulse candidates on plots.

singlepulse_maxwidth The maximum pulse width (in seconds) for singlepulse searching.

to_prepfold_sigma The minimum incoherent sum significance required to fold a candidate.

max_cands_to_fold The maximum number of candidates to fold per beam.

numhits_to_fold The minimum number of DMs at which a periodicity must be found to be folded.

low_DM_cutoff The lowest DM to consider when choosing candidates to fold.

lo_accel_numharm The maximum number of harmonics to sum in low acceleration searching.

lo_accel_sigma The minimum gaussian significance threshold in low acceleration searching.

lo_accel_zmax The maximum z-value (in bins) used in low acceleration searching.

lo_accel_flo The maximum frequency searched in low acceleration searching.

hi_accel_numharm The maximum number of harmonics to sum in high acceleration searching.

hi_accel_sigma The minimum gaussian significance threshold in high acceleration searching.

hi_accel_zmax The maximum z-value (in bins) used in high acceleration searching.

hi_accel_flo The maximum frequency searched in high acceleration searching.

low_T_to_search The shortest observation that will be searched (in seconds).

base_tmp_dir The path where the temporary working directory should be created. This could be /dev/shm, for example. If /dev/shm doesn't exist, or is too small, this path could simply be another path in the scratch area of the worker node.

sifting_sigma_threshold The minimum incoherent power sigma to consider when sifting candidates. This should be left at `to_prepfold_sigma-1.0`.

sifting_c_pow_threshold The coherent power threshold used when sifting.

sifting_r_err The Fourier bin tolerance to consider two candidates as being at the same frequency.

sifting_short_period The shortest candidate period to consider (in seconds).

sifting_long_period The longest candidate period to consider (in seconds).

sifting_harm_pow_cutoff The power required in at least one harmonic.

3.5.9 upload.py

Configurations used by the uploader are set in `upload.py`.

version_num The version number of the current version of the pipeline (including versions of `PRESTO` and `psrfits_utils`). The preferred format is:

```
PRESTO:{git-hash};PIPELINE:{git-hash};PSRFITS_UTILS{git-hash}.
```

It is **strongly** recommended that the git-hash values are retrieved dynamically, so they are always up-to date. However, if this is not possible you are able to hard code values of the git-hashes, but they should be changed every time the versions of the software used are modified.

3.6 Job-Tracker Database Details

The job-tracker database contains six tables: `requests`, `files`, `download_attempts`, `jobs`, `job_files`, and `job_submits`. Each of these tables is meant to track a particular object (e.g. `jobs`), or events (e.g. job submissions to the queue).

A description of each database table will now be presented:

3.6.1 `requests`

The `requests` table tracks requests for data restores. An entry in the table is inserted when a request is made. Each entry contains the following data:

id (UNIQUE INTEGER) A unique identifier, automatically assigned by the database upon creation of an entry.

size (INTEGER) Total size of the files restored, in bytes.

guid (TEXT) A globally unique identifier for the restore. This serves as the directory name on the FTP server where the restored files will be placed.

status (TEXT) The current status of the request.

details (TEXT) Details about the request.

created_at (TEXT) Date and time when the request entry was created.

updated_at (TEXT) Date and time when the request entry was last updated.

3.6.2 `files`

The `files` table tracks each file available to the pipeline. A separate entry is inserted for each file. Each entry contains the following data:

id (UNIQUE INTEGER) A unique identifier, automatically assigned by the database upon creation of an entry.

filename (TEXT) The name of the file.

remote_filename (TEXT) The name and path of the downloaded file on the FTP server.

request_id (INTEGER) The identifier of the entry in the **requests** table that restored this file.

status (TEXT) The current status of the file.

details (TEXT) Details about the file.

created_at (TEXT) Date and time when the file entry was created.

updated_at (TEXT) Date and time when the file entry was last updated.

3.6.3 download_attempts

Every time one an attempt to download one of the files listed in the **files** table an entry is created in the **download_attempts** table. Each download attempt entry has the following data:

id (UNIQUE INTEGER) A unique identifier, automatically assigned by the database upon creation of an entry.

request_id (INTEGER) The identifier of the entry in the **files** table that this download attempt corresponds to.

status (TEXT) The current status of the download attempt.

details (TEXT) Details about the current status of the download attempt.

created_at (TEXT) Date and time when the download attempt entry was created.

updated_at (TEXT) Date and time when the download attempt entry was last updated.

3.6.4 jobs

Each beam-worth of data gets submitted to be processed as a single job. These jobs are tracked in the `jobs` table, and have the following data:

id (UNIQUE INTEGER) A unique identifier, automatically assigned by the database upon creation of an entry.

status (TEXT) The current status of the job.

details (TEXT) Details about the current status of the job.

created_at (TEXT) Date and time when the job entry was created.

updated_at (TEXT) Date and time when the job entry was last updated.

3.6.5 job_files

Because a job can have multiple files a mapping between jobs and files is required. The `job_files` table serves this purpose. Each row in the table has the following data:

id (UNIQUE INTEGER) A unique identifier, automatically assigned by the database upon creation of an entry.

file_id (INTEGER) The identifier of an entry in the `files` table.

job_id (INTEGER) The identifier of an entry in the `jobs` table that the file referred to by `file_id` belongs to.

created_at (TEXT) Date and time when the entry was created.

updated_at (TEXT) Date and time when the entry was last updated.

3.6.6 job_submits

Each submission of a job to the queue to be processed is tracked in the `job_submits` table. Each job submission has the following information:

id (UNIQUE INTEGER) A unique identifier, automatically assigned by the database upon creation of an entry.

queue_id (TEXT) The unique identifier assigned to the job by the queue manager when it is submitted. This identifier is used to check if the submission is still being processed.

job_id (INTEGER) The identifier of the job (from the `jobs` table) that was submitted for processing.

output_dir (TEXT) The directory where the results of processing will be output.

status (TEXT) The current status of the job.

details (TEXT) Details about the current status of the job.

created_at (TEXT) Date and time when the entry was created.

updated_at (TEXT) Date and time when the entry was last updated.